

The image shows a screenshot of a Mettl online test interface. At the top, there's a navigation bar with the Mettl logo, 'Section 1 of 1', 'Section #1', and a 'Revisit' button. Below the navigation, a blue header bar contains the text 'Question #1' and 'Is container orchestration system'. To the right of this, a sidebar titled 'Choose the best option' lists four options: 'Pods', 'Ranger', 'Docker', and 'Kubernetes'. A small mouse cursor icon is positioned near the bottom of the sidebar. The main content area below the header is mostly blank.

mettl

Section 1 of 1 Section #1

Revisit

Question #1

Is container orchestration system

Choose the best option

- Pods
- Ranger
- Docker
- Kubernetes

Kubernetes

narzo Shot by meghs

+91-08047189190

Question # 2

Which is not true for Pig?

Choose the best option

- It is a scripting language
- It is necessary to define a schema
- It uses grunt shell
- It runs MR jobs under the hood

It runs MR jobs under the hood

Section 1 of 1 Section #1

Revisit

Question #3

What is the command to check Hadoop version?

hadoop version

Choose the best option

- hadoop -version
- hadoop version
- hadoop -version
- hadoop -v

Question #4

Optimizations can be performed while using MR jobs by using a combiner to work on the mapper output before it is sent to the reducer and by increasing number of reducers

Choose the best option

- True
- False

True

Question # 5

Which is the default replication factor for HDFS?

- 1
- 2
- 3
- 0

3

Question # 6

Which file(s) is/are used for storing meta data in Name node?

- editlogs
- fsimage
- Both A & B
- None of the above

Both A&B

Question # 7

From the below steps, choose right sequence of steps required when an HDFS Client wants to create a File on the Data Node ?

- 1) HDFS Client Write the Packet on the Data Node.
- 2) HDFS Client gets the block location from the Name Node.
- 3) Data Node will send acknowledgement back to the HDFS Client.
- 4) HDFS Client will send a create file request to the Name Node.

Choose the most appropriate Sequence from the below options.

Choose the best option

- 1 2 3 4
- 2 1 3 4
- 4 2 1 3
- 2 3 4 1

Question # 7

From the below steps, choose right sequence of steps required when an HDFS Client wants to create a File on the Data Node ?

- 1) HDFS Client Write the Packet on the Data Node.
- 2) HDFS Client gets the block location from the Name Node.
- 3) Data Node will send acknowledgement back to the HDFS Client.
- 4) HDFS Client will send a create file request to the Name Node.

Choose the most appropriate Sequence from the below options.

Choose the best option

- 1 2 3 4
- 2 1 3 4
- 4 2 1 3
- 2 3 4 1

Question # 8

Which of the following phases occur simultaneously in Map-Reduce Framework?
Choose the most appropriate option.

- Reduce and Sort
- Shuffle and Sort
- Shuffle and Map
- All of the above

Shuffle & Sort

Question # 8

Which file contains configuration for resource manager?

Choose the correct answer

- yarn-site.xml
- mapred-site.xml
- core-site.xml
- hdfs-site.xml

yarn-site.xml

Question # 10

Which of the following process schedules the tasks across nodes in Hadoop system



- Node manager
- Resource manager
- scheduler
- Application master

Resource Manager

Question # 11

Which datatype can be used to implement a bounded variable length character string?

- String
- char
- varchar
- none of the above

String (not sure)

Question # 12

What does the Hive metastore do?

- runs the Hiveql queries
- exposes the HDFS in form of relational tables
- executes MR jobs on Hive
- none of the above

none of the above

Question # 13

Pig operates in mainly how many modes?

2

Revisit

Choose the correct answer

- 2
- 3
- 4
- 5

Question # 14

All MapReduce jobs reading from an HBase table accept their [K1, V1] pair in the form of

rowkey:scan result

- rowid:cell value
- rowkey:scan result
- column Family:cell value
- column attribute:scan result

Question # 15

Assume that the Spark Context reference is available as "sc", which of the following statement will help us create a Spark SQL Context "ssc" successfully ?

- ssc = new SQLContext(sc)
- ssc = SQLContext(sc)
- ssc = sc.SQLContext()
- ssc = SQLContext()

ssc=SQLContext(sc)

Question # 16

If we have to create a Structure or schema for a Semi-Structured data file having the below fields

Product_ID Integer
Product_Name String
Seller_ID Integer
Seller_Name String

Which of the following code can help us do the same ?

A (not sure)

Choose the best option

- Schema1 = StructType({StructField("Product_ID", IntegerType(), True),
StructField("Product_Name", StringType(), True),
StructField("Seller_ID", IntegerType(), True),
StructField("Seller_Name", StringType(), True)})
- Schema1 = StructType({("Product_ID", IntegerType(), True),
("Product_Name", StringType(), True),
("Seller_ID", IntegerType(), True),
("Seller_Name", StringType(), True)})
- Schema1 = StructField({("Product_ID", IntegerType(), True),
("Product_Name", StringType(), True),
("Seller_ID", IntegerType(), True),
("Seller_Name", StringType(), True)})
- Schema1 = StructFields({StructField("Product_ID", IntegerType(), True),
StructField("Product_Name", StringType(), True),
StructField("Seller_ID", IntegerType(), True),
StructField("Seller_Name", StringType(), True)})

Question # 17

Which attribute is not indicative for data streaming?

Choose the best option

- Limited processing power
- Limited memory
- Limited data
- Limited time

Limited Data

Question # 18

Which method/s are used while developing windowing based streaming application? Select All that apply (Check Box)

- reduceByKey
- countByValueAndWindow
- groupByKey
- reduceByKeyAndWindow

Question # 19

Which co-ordination service is used in Kafka Cluster -

- Impala
- ActiveMQ
- BigTop
- Zookeeper

ZooKeeper

Question # 20

Which of the following statement is True about Apache Kafka ?

- Each Kafka topic has one or more partitions
- Each partition is maintained on at least one or more Brokers
- both option 1 and 2

Both 1 & 2

Question # 20

Which of the following statement is True about Apache Kafka?

Choose the best option

- Each Kafka topic has one or more partitions
- Each partition is maintained on at least one or more Brokers
- both option 1 and 2